

Azure Data Factory

Lab: Orchestrating Data Movement with Azure Data Factory

Pre-requisites:

- Azure Pass subscription
- Azure Data Lake Storage Gen2 storage account.
- Azure SQL Database

Lab objectives:

1. Ingest data using the Copy Activity
2. Use the Mapping Data Flow task to perform transformation

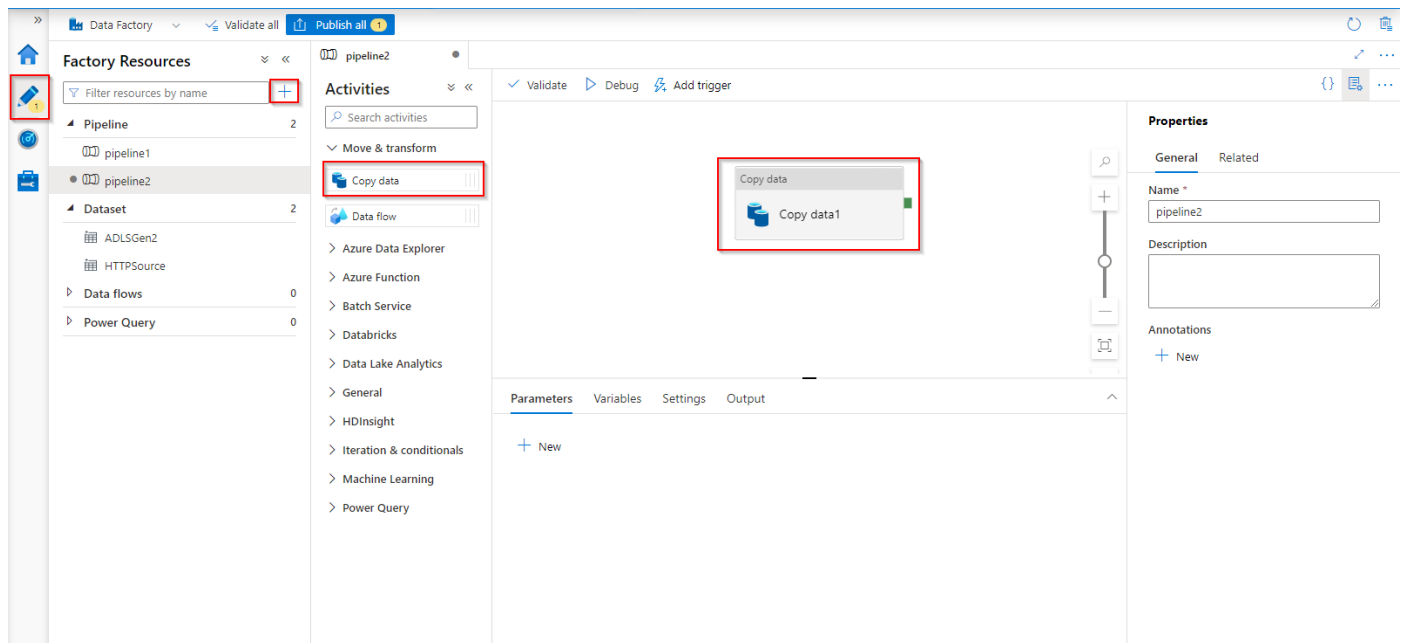
Exercise 1: Use the Mapping Data Flow task to perform transformation

The task for this exercise are as follows

1. Add the Copy Activity to the pipeline
2. Create a new HTTP dataset to use as a source
3. Create a new ADLS Gen2 sink
4. Test the Copy Activity

Task 1: Add the Copy Activity to the designer

1. Launch Data Factory Studio from Azure Portal.
2. **Open the authoring canvas** If coming from the ADF homepage, click on the **pencil icon** on the left sidebar and select the **+ pipeline button** to open the authoring canvas and create a pipeline.
3. **Add a copy activity** In the Activities pane, open the **Move and Transform** accordion and drag the **Copy data** activity onto the pipeline canvas.




Task 2: Create a new HTTP dataset to use as a source

1. In the Source tab of the Copy activity settings, click **+ New**
2. In the data store list, select the **HTTP** tile and click **continue**
3. In the file format list, select the **DelimitedText** format tile and click **continue**
4. In **Set properties** blade, give your dataset an understandable name such as **HTTPSource** and click on the **Linked Service** dropdown. If you have not created your HTTP Linked Service, select **New**.
5. In the **New Linked Service (HTTP)** screen, copy the URL of the moviesDB csv file below in the **Base URL** textbox. You can access the data with no authentication required using the following endpoint:

<https://raw.githubusercontent.com/djpmstf/adf-ready-demo/master/moviesDB.csv>

New linked service

 HTTP [Learn more](#)

Name *

HttpServer2

Description

Connect via integration runtime * ⓘ

AutoResolveIntegrationRuntime

Base URL *

https://raw.githubusercontent.com/djpmstf/adf-ready-demo/master/moviesDB.csv

Server Certificate Validation ⓘ

☒ Enable ☐ Disable

Authentication type *

Anonymous

Auth headers ⓘ

[+ New](#)

Annotations


[+ New](#)

[> Parameters](#)

[> Advanced ⓘ](#)

Create

Cancel

 Test connection

6. In the **Authentication type** drop down, select **Anonymous**. and click on **Create**.

- Once you have created and selected the linked service, specify the rest of your dataset settings. These settings specify how and where in your connection we want to pull the data. As the url is pointed at the file already, no relative endpoint is required. As the data has a header in the first row, set **First row as header** to be true and select Import schema from **connection/store** to pull the schema from the file itself. Select **Get** as the request method. You will see the following screen

Set properties

Name

HttpSourceMovie

Linked service *

HttpServer2

Relative URL

First row as header



Import schema



From connection/store



From sample file



None

> Advanced

OK

Back

Cancel

- Click **OK** once completed.

To verify your dataset is configured correctly, click **Preview data** in the Source tab of the copy activity to get a small snapshot of your data.

Preview data

Linked service: HttpServer2

Object:

movie	title	genres	year	Rating	Rotton Tomato
108583	Fawlty Towers (1975	Comedy	-1980	1	54
32898	Trip to the Moon, A (Voyage dans la lune, Le)	Action Adventure Fantasy Sci-Fi	1902	7	80
7065	Birth of a Nation, The	Drama War	1915	6	92
7243	Intolerance: Love's Struggle Throughout the Ages	Drama	1915	4	82

Task 3: Create a new ADLS Gen2 dataset sink

1. Click on the **Sink tab**, and then click **+ New**
2. Select the **Azure Data Lake Storage Gen2** tile and click **Continue**.
3. Select the **DelimitedText** format tile and click **Continue**.
4. In the **Set Properties** blade, give your dataset an understandable name such as **ADLSGen2Movie** and click on the **Linked Service** dropdown. If you have not created your ADLS Linked Service, select **New**.
5. In the **New linked service (Azure Data Lake Storage Gen2)** blade, select your authentication method as **Account key**, select your **Azure Subscription** and select your Storage account name of **datalakexx**. You will see a screen as follows:
6. Click on **Create**
7. Once you have configured your linked service, you enter the **set properties** blade. As you are writing to this dataset, you want to point the folder where you want moviesDB.csv copied to. In the example below, I am writing to folder **output** in the **data** container. Set **First row as header** to be true and Import schema set to **None**.

Set properties

Name

ADLSGen2Movie

Linked service *

ADLSGen2

File path

data

/

output

/

moviesDB.csv



First row as header



Import schema

☐ From connection/store

☐ From sample file

☒ None

> Advanced

OK

Back

Cancel

- Click **OK** once completed.

Task 4: Test the Copy Activity

At this point, you have fully configured your copy activity. To test it out, click on the **Debug** button at the top of the pipeline canvas. This will start a pipeline debug run.

1. To monitor the progress of a pipeline debug run, click on the **Output** tab of the pipeline
2. To view a more detailed description of the activity output, click on the eyeglasses icon. This will open up the copy monitoring screen which provides useful metrics such as Data read/written, throughput and in-depth duration statistics.

The screenshot shows the Azure Data Factory pipeline canvas. At the top, there are three buttons: 'Validate' (with a checkmark icon), 'Debug' (with a play icon, highlighted with a red box), and 'Add trigger' (with a lightning bolt icon). Below these buttons is a canvas area containing a single activity named 'Copy data1', which is represented by a blue cube icon and a green checkmark. To the right of the canvas is a vertical toolbar with icons for search, add, zoom in, zoom out, and full screen. Below the canvas is a tabbed interface with four tabs: 'Parameters', 'Variables', 'Settings', and 'Output' (highlighted with a red box). Below the tabs is a section for the pipeline run, showing the 'Pipeline run ID: 13eb76cf-42b7-4661-b288-bffbc22532a' and a 'View debug run consumption' link. Below this is a table with the following columns: Name, Type, Run start, Duration, and Status.

Name	Type	Run start	Duration	Status
Copy data1	Copy data	2022-03-01T13:26:24.721	00:00:09	Success

3. To verify the copy worked as expected, open up your ADLS gen2 storage account and check to see your file was written as expected

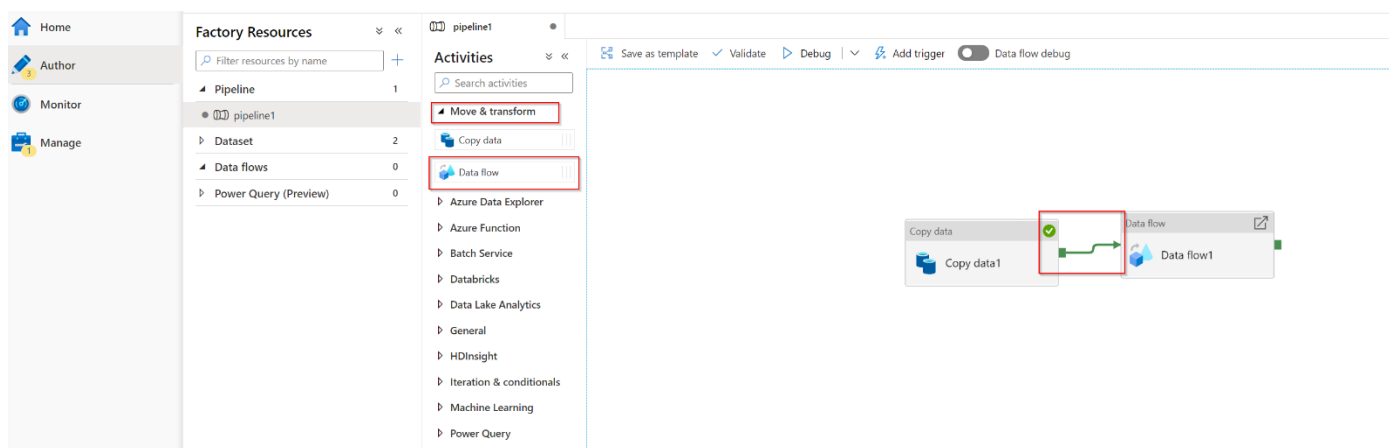
Exercise 2: Transforming Data with Mapping Data Flow

Tasks for this exercise are as follows:

1. Preparing the environment
2. Adding a Data Source
3. Using Mapping Data Flow transformation
4. Writing to a Data Sink
5. Running the Pipeline

Task 1: Preparing the environment

1. **Add a Data Flow activity** In the Activities pane, open the **Move and Transform** accordion and drag the **Data Flow** activity onto the pipeline canvas and connect with the **Copy data** activity we created earlier.



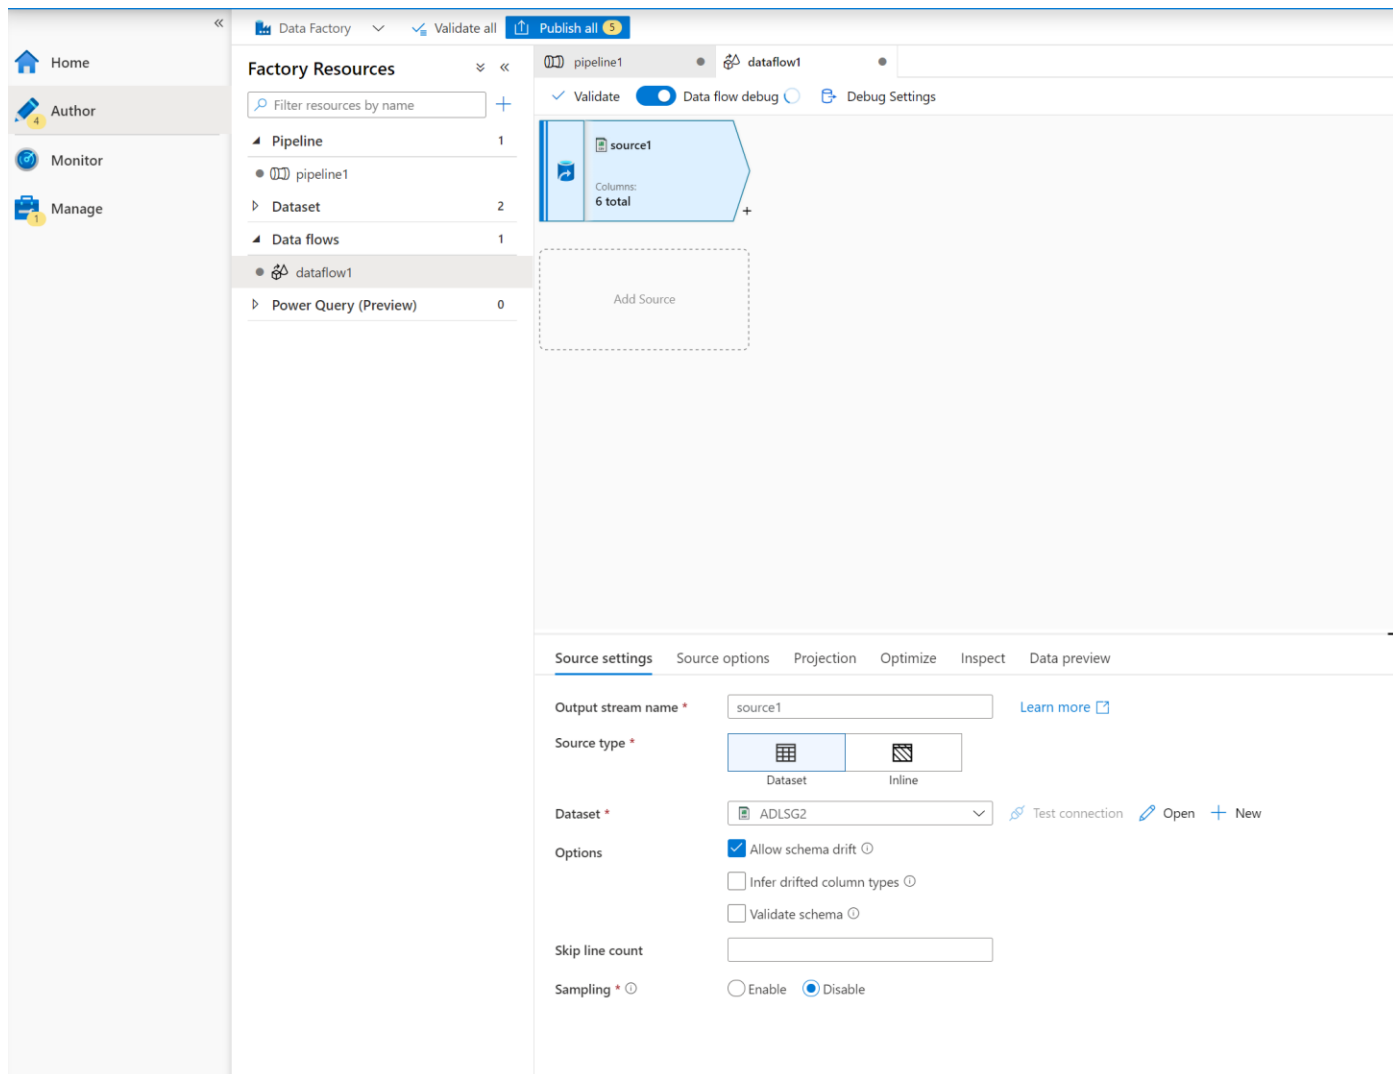
- 2.
3. Turn the **Data Flow Debug** slider located at the top of the authoring module on, and click **OK** in the **Turn on data flow debug** screen that appears.

NOTE: Data Flow clusters take 5-7 minutes to warm up.

4. Select the data flow activity in the pipeline workspace. In the lower pane, select the settings tab, click **+ New** for the variable **Dataflow**

Task 2: Adding a Data Source

1. **Add an ADLS source:** Click on the Mapping Data Flow object in the canvas. Go to the source settings tab. In the **Dataset** dropdown, select your **ADLSGen2** dataset used in your Copy activity

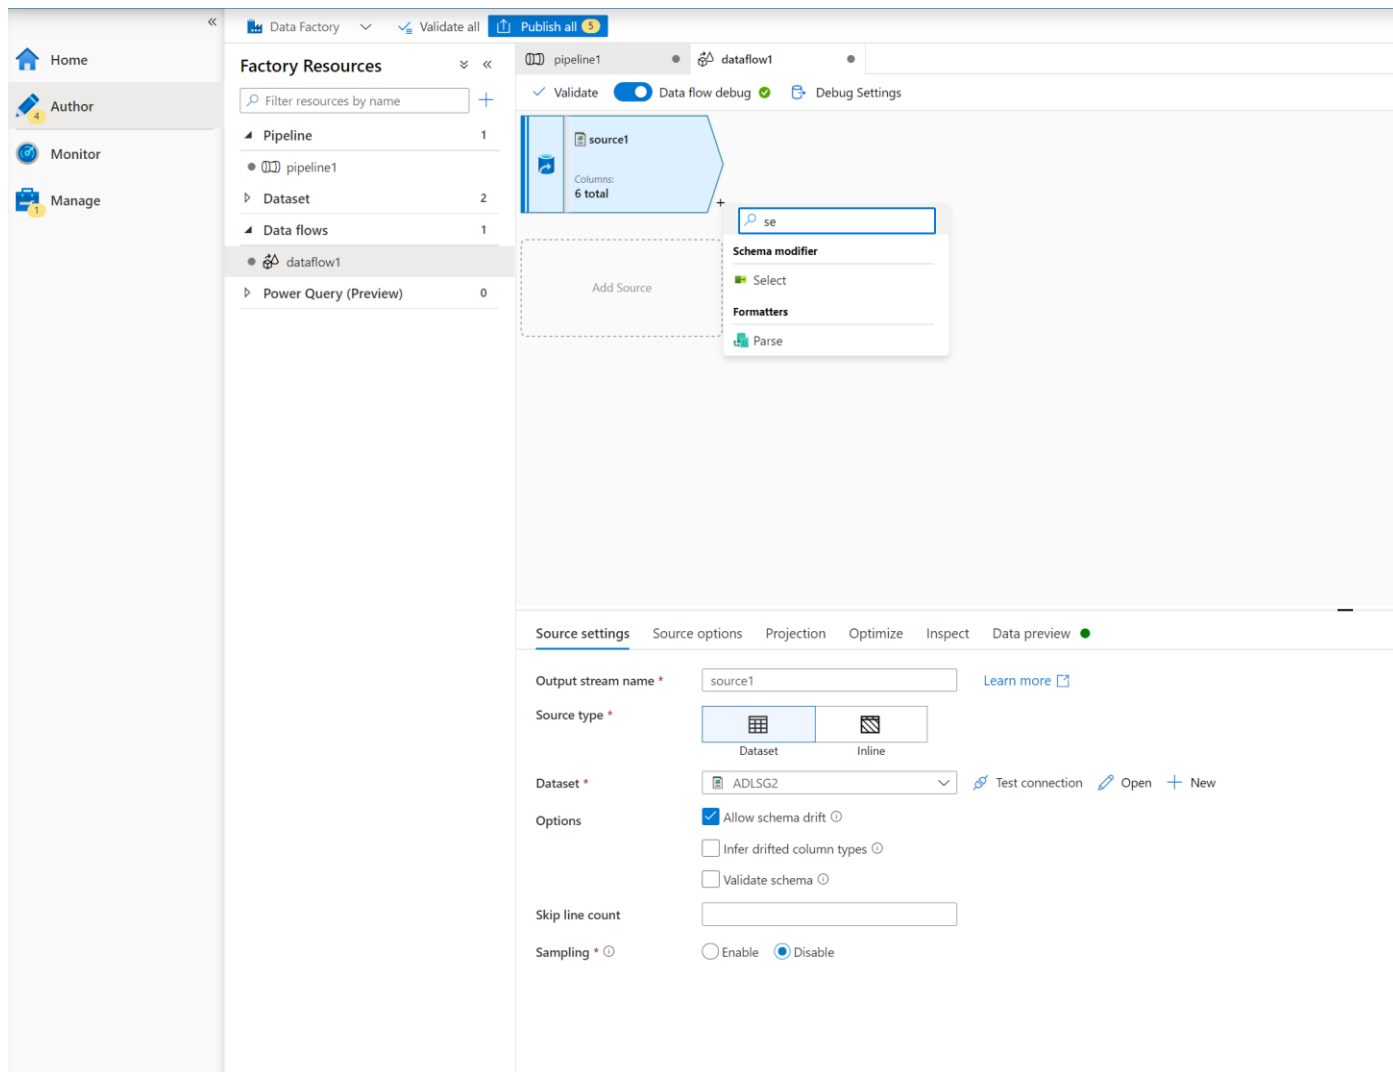


- If your dataset is pointing at a folder with other files, you may need to create another dataset or utilize parameterization to make sure only the moviesDB.csv file is read
- If you have not imported your schema in your ADLS, but have already ingested your data, go to the dataset's 'Schema' tab and click 'Import schema' so that your data flow knows the schema projection.

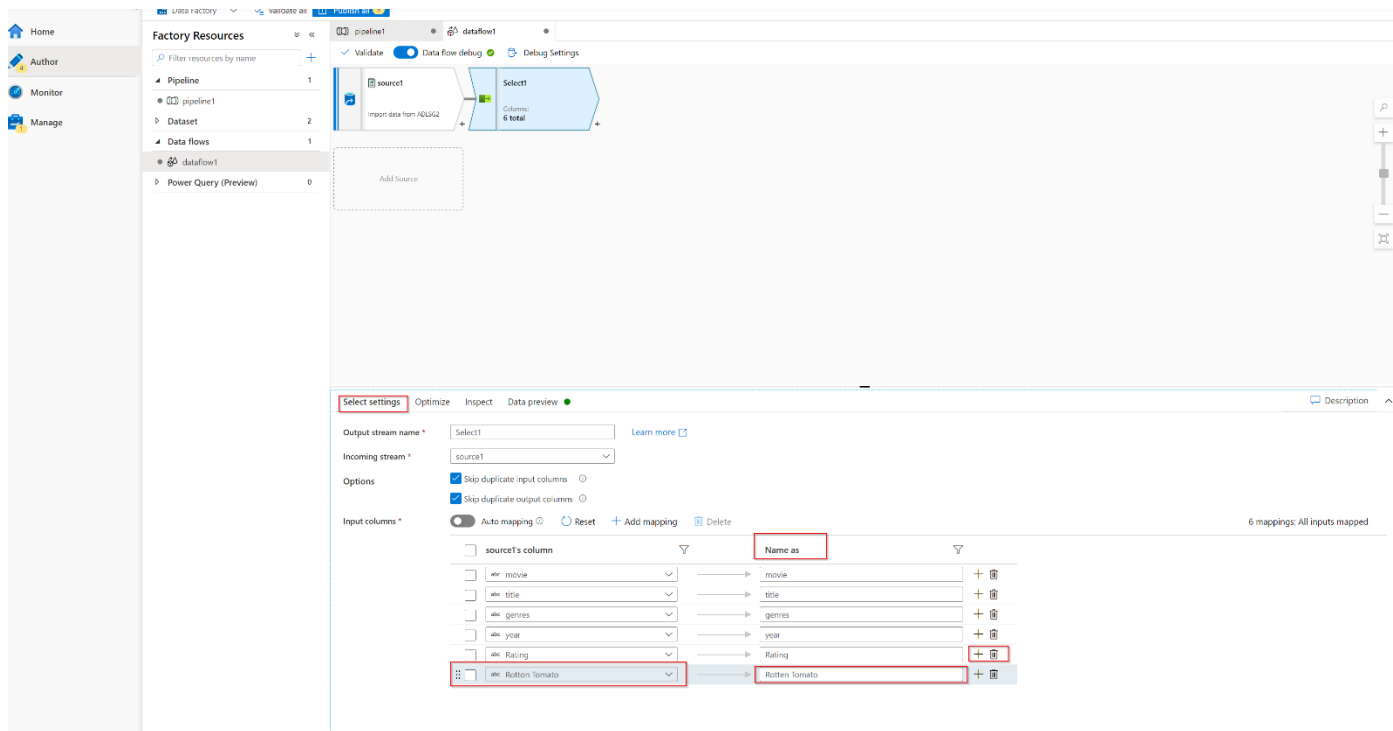
Once your debug cluster is warmed up, verify your data is loaded correctly via the **Data preview** tab. Once you click the refresh button, Mapping Data Flow will show calculate a snapshot of what your data looks like when it is at each transformation.

Task 3: Using Mapping Data Flow transformation

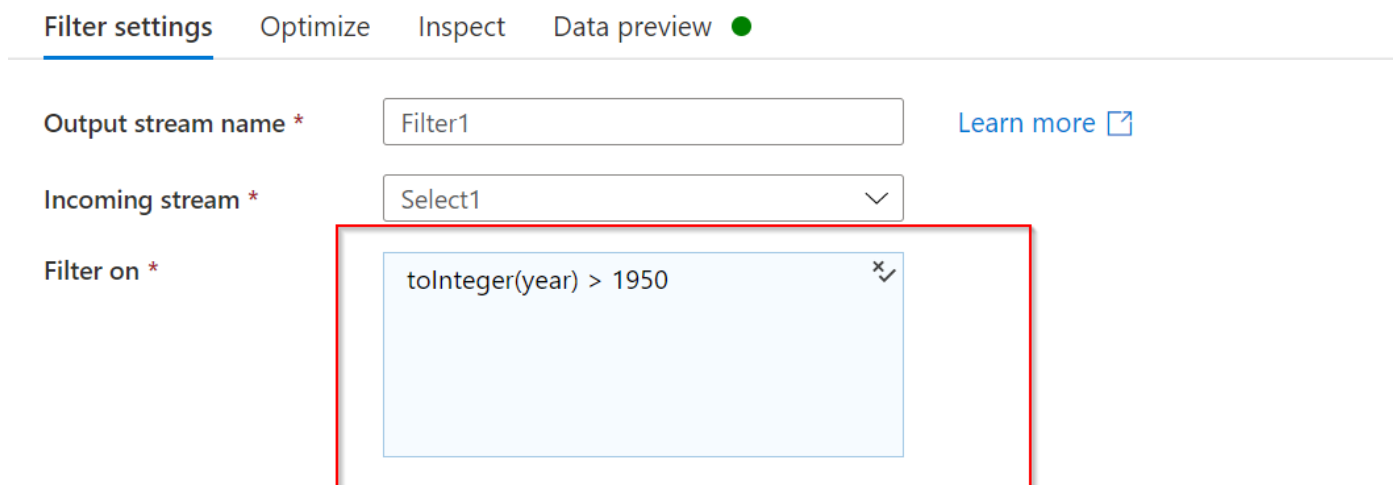
1. **Add a Select transformation to rename and drop a column:** In the preview of the data, you may have noticed that the "Rotton Tomatoes" column is misspelled. To correctly name it and drop the unused Rating column, you can add a Select transformation by clicking on the + icon next to your ADLS source node and choosing **Select** under Schema modifier.







In the **Name as** field, under the **Select settings** tab, change 'Rotton' to 'Rotten'. To drop the Rating column, hover over it and click on the trash can icon.



2. **Add a Filter Transformation to filter out unwanted years:** Say you are only interested in movies made after 1951. You can add a Filter transformation to specify a filter condition by clicking on the **+** icon next to your Select transformation and choosing **Filter** under Row Modifier. Click on the **expression box** to open up the Expression builder and enter in your filter condition. Using the syntax of the Mapping Data Flow expression language, **toInteger(year) > 1950** will convert the string year value to an integer and filter rows if that value is above 1950.



When you clicked on **open expression builder** you can verify your condition is working properly. This will also show by a check mark in the **Filter on** textbox.



Visual expression builder

Filter1

Expression

toInteger(year) > 1950

+

-

*

/

||

&&

!

Expression elements

All

Functions

Input schema

Parameters

Cached lookup

Expression values

Filter by keyword

+ Create new

abc movie

abc title

abc genres

abc year

abc Rotten Tomato

123 abs(123 numeric_value)

123 acos(123 numeric_value)

ANY add(ANY first_expression, ANY second_expression)

addDays(date/timestamp, ANY days to add)

addMonths(date/timestamp, ANY months to add)

and(x first_condition, x second_condition)

[] array(ANY items)

Data preview

Refresh

Save and finish

Cancel

Clear contents

3. **Add a Derive Transformation to calculate primary genre:** As you may have noticed, the genres column is a string delimited by a '|' character. If you only care about

the *first* genre in each column, you can derive a new column named **PrimaryGenre** via the Derived Column transformation by clicking on the **+** icon next to your Filter transformation and choosing **Derived Column** under Schema Modifier. Similar to the filter transformation, the derived column uses the Mapping Data Flow expression builder to specify the values of the new column.

The screenshot displays the Data Studio interface. At the top, a pipeline is visible with the following steps: **source1** (Import data from ADLSG2), **Select1** (Renaming source1 to Select1 with columns 'movie, title, genres, year, Rotten Tomato'), **Filter1** (Filtering rows using expressions on columns 'year'), and **DerivedColumn1** (Columns: 6 total). Below the pipeline is an "Add Source" button. The bottom panel shows the **Derived column's settings** tab, which includes fields for **Output stream name *** (DerivedColumn1) and **Incoming stream *** (Filter1). Under the **Columns *** section, there is a table with two columns: **Column** and **Expression**. The first row shows **PrimaryGenre** in the **Column** column and the expression `iif(locate(' ', genres) > 1, left(genres, locate(' ', ge... abc` in the **Expression** column. The expression is highlighted with a red box.

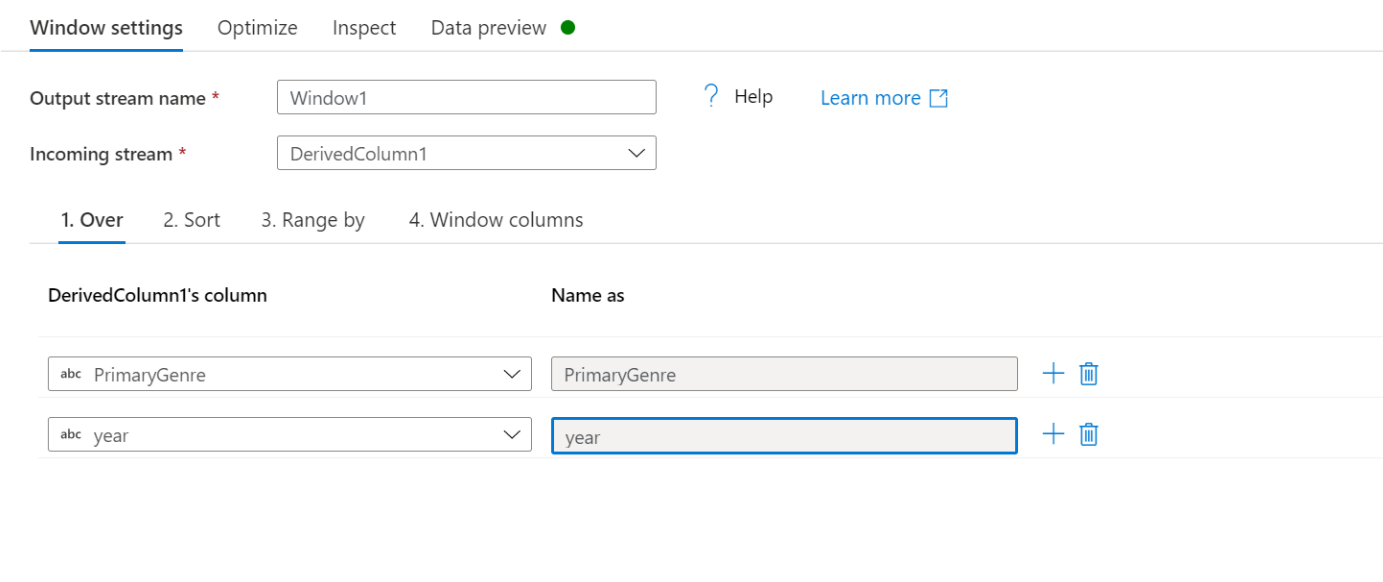
In this scenario, you are trying to extract the first genre from the genres column which is formatted as 'genre1|genre2|...|genreN'. Use the **locate** function to get the first 1-based index of the '|' in the genres string. Using the **iif** function, if this index is greater than 1, the primary genre can be calculated via the **left** function which returns all characters in a string to the left of an index. Otherwise, the PrimaryGenre value is equal to the genres field. You can verify the output via the expression builder's Data preview pane.

- In the **Derived column's** settings tab, click **+Add** then, **Add column**, to add a column named **PrimaryGenre**.
- Under **Expression** open the **Expression builder**.

- Write `iif(locate('|', genres)>1,left(genres,locate('|', genres)-1),genres)`
- Select **Save and finish**

4. **Rank movies via a Window Transformation** Say you are interested in how a movie ranks within its year for its specific genre. You can add a Window transformation to define window-based aggregations by clicking on the **+** icon next to your Derived Column transformation and clicking **Window** under Schema modifier. To accomplish this, specify what you are windowing over, what you are sorting by, what the range is, and how to calculate your new window columns. In this example, we will window over PrimaryGenre and year with an unbounded range, sort by Rotten Tomato descending, and calculate a new column called RatingsRank which is equal to the rank each movie has within its specific genre-year.

- In the **Window settings** pane under the **Over** tab, select **PrimaryGenre** and add **year** by clicking on **+** and selecting **year** from the dropdown.



Window settings Optimize Inspect Data preview ●

Output stream name * ? Help [Learn more](#)

Incoming stream *

1. Over 2. Sort 3. Range by 4. Window columns

DerivedColumn1's column	Name as
abc PrimaryGenre	PrimaryGenre
abc year	year

- In the **Sort settings** pane, select the **Rotten Tomato** column, select **Descending** under **Order** and check **Nulls first**

Window settings Optimize Inspect Data preview ●

Output stream name *

Window1

? Help

[Learn more](#)

Incoming stream *

DerivedColumn1



1. Over

2. Sort

3. Range by

4. Window columns

DerivedColumn1's column

Order

Nulls first

abc Rotten Tomato



Descending



- In the **Range by settings** pane, leave all settings per default.

Window settings Optimize Inspect Data preview ●

Output stream name *

Window1

? Help

[Learn more](#)

Incoming stream *

DerivedColumn1



1. Over

2. Sort

3. Range by

4. Window columns

Option * ⓘ



Range by current row offset



Range by column value

Unbounded



- In the **Window columns settings** pane, rename the blank column to **RatingsRank** and enter as expression **rank()**

Window settings
Optimize
Inspect
Data preview

Output stream name *
Window1
? Help
Learn more

Incoming stream *
DerivedColumn1

1. Over
2. Sort
3. Range by
4. Window columns

+ Add
Clone
Delete
Open expression builder

Column	Expression
<input type="checkbox"/> RatingsRank	rank() 123 +

5. **Aggregate ratings with an Aggregate Transformation:** Now that you have gathered and derived all your required data, we can add an Aggregate transformation to calculate metrics based on a desired group by clicking on the **+** icon next to your Window transformation and clicking **Aggregate** under Schema modifier. As you did in the window transformation, lets group movies by PrimaryGenre and year
- Under the **Aggregate settings** tab, select **Group by**.
 - Using the dropdown select the column **Primary Genre** and add the **year** column by clicking **+**, and dropdown.

Aggregate settings
Optimize
Inspect
Data preview

Output stream name *
Aggregate1
Learn more

Incoming stream *
Window1

Group by
Aggregates

Columns	Name as
abc PrimaryGenre	PrimaryGenre +
abc year	year +

In the Aggregates tab, you can aggregations calculated over the specified group by columns. For every genre and year, lets get the average Rotten Tomatoes rating, the highest and lowest rated movie (utilizing the windowing function) and the number of movies that are in each group. Aggregation significantly reduces the amount of rows in your transformation stream and only propagates the group by and aggregate columns specified in the transformation.

- Under the **Aggregate settings** tab, now select **Aggregates**. Add the following columns by clicking **+** and then **Add column**, with their respective expressions:
 - AverageRating: avg(toInteger({Rotten Tomato}))
 - HighestRead: first(title)
 - LowestRead: last(title)
 - NumberOfMovies: count()

Aggregate settings
 Optimize
 Inspect
 Data preview

Output stream name * [Learn more](#)

Incoming stream *

Group by Aggregates

Grouped by: PrimaryGenre, year

[+](#) Add
 [📄](#) Clone
 [🗑️](#) Delete
 [🔗](#) Open expression builder

<input type="checkbox"/>	Column	Expression
<input type="checkbox"/>	AverageRating	avg(toInteger({Rotten Tomato})) 1.2 + 🗑️
<input type="checkbox"/>	HighestRead	first(title) abc + 🗑️
<input type="checkbox"/>	LowestRead	last(title) abc + 🗑️
<input type="checkbox"/>	NumberOfMovies	count() 12L + 🗑️

- To see how the aggregate transformation changes your data, use the Data Preview tab

6. **Specify Upsert condition via an Alter Row Transformation** If you are writing to a tabular sink, you can specify insert, delete, update and upsert policies on rows using the Alter Row transformation by clicking on the **+** icon next to your Aggregate transformation and clicking **Alter Row** under Row modifier. Since you are always inserting and updating, you can specify that all rows will always be upserted.

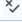
- From the dropdown next to **Alter row conditions** in the **Alter row settings** tab, please select **Upsert if**. In the expression write **true()**

Output stream name * [Learn more](#) 

Incoming stream * 

Alter row conditions * 

 Upsert if 







Task 4: Writing to a Data Sink

1. **Write to a Azure SQL Database Sink:** Now that you have finished all your transformation logic, you are ready to write to a Sink.
 - i. Add a **Sink** by clicking on the **+ icon** next to your Alter row transformation and clicking **Sink** under Destination.
 - ii. In the Sink tab, create a new data warehouse dataset via the **+ New button** next to **Dataset**.
 - iii. Select **Azure SQL Database** from the tile list and click **Continue**
 - iv. Select **+New** under **Linked service**. Configure your Azure SQL Database connection to connect to the SQLDB database.
 - v. **Account selection method: From Azure subscription**
 - vi. **Azure subscription:** select the subscription used for this lab.
 - vii. **Server name:** select your **sqlservicexx** server.
 - viii. **Database name: SQLDB**
 - ix. **Authentication type: SQL authentication**
 - x. For **username** use your server admin username, for **Password** use the corresponding password you provided, when setting up the service.


New linked service

 Azure SQL Database [Learn more](#) 

Name *

AzureSqlDatabase

Description

Connect via integration runtime * 

AutoResolveIntegrationRuntime

Connection string

Azure Key Vault

Account selection method 

☒ From Azure subscription ☐ Enter manually

Azure subscription

Azure Pass - Sponsorship (c8fd5e96-3329-4700-b2ef-455eb9a984e4)

Server name *

sqlservicepc

Database name *

SQLDB

Authentication type *


SQL authentication

User name *

sqladmin

Create

Cancel

 Test connection

- xi. Click **Create** when finished.
- xii. In the **Set properties** page, select **Create new table** and enter in the schema of **dbo** and the table name of **Ratings**. Click **OK** once completed.

Set properties

Name

AzureSqlTable

Linked service *

AzureSqlDatabase

☐ Select from existing table ☒ Create new table

Schema and table name

dbo

Ratings

> Advanced

- xiii. Since an upsert condition was specified, you need to go to the Settings tab and select **Allow upsert**.

✓ Validate ☒ Data flow debug ☒ Debug Settings

Sink **Settings** Mapping Optimize Inspect Data preview

Update method ^①

☐ Allow insert

☐ Allow delete

☒ Allow upsert

☐ Allow update

Key columns ^① ☒ List of columns ☐ Custom expression ^①

abc PrimaryGenre +

abc year +

Skip writing key columns ☐

Table action ^① ☒ None ☐ Recreate table ^① ☐ Truncate table ^①

Batch size ^①

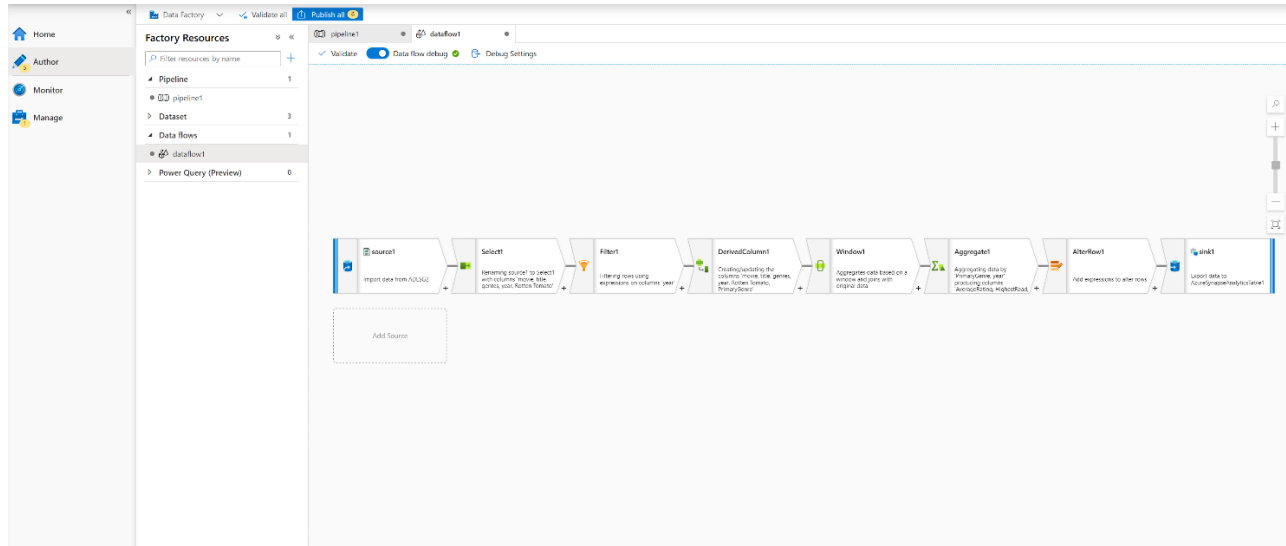
Use TempDB ^① ☒

Pre SQL scripts ^① ☒ List of scripts ☐ Custom expression ^①

+


- xiv. For **Key columns** select **List of Columns** and add through **+** the two columns PrimaryGenre and year. based on key columns PrimaryGenre and year.
- xv. In the **Mapping** pane make sure you untick **Auto mapping**.

At this point, You have finished building your 8 transformation Mapping Data Flow. It's time to publish and run the pipeline and see the results!

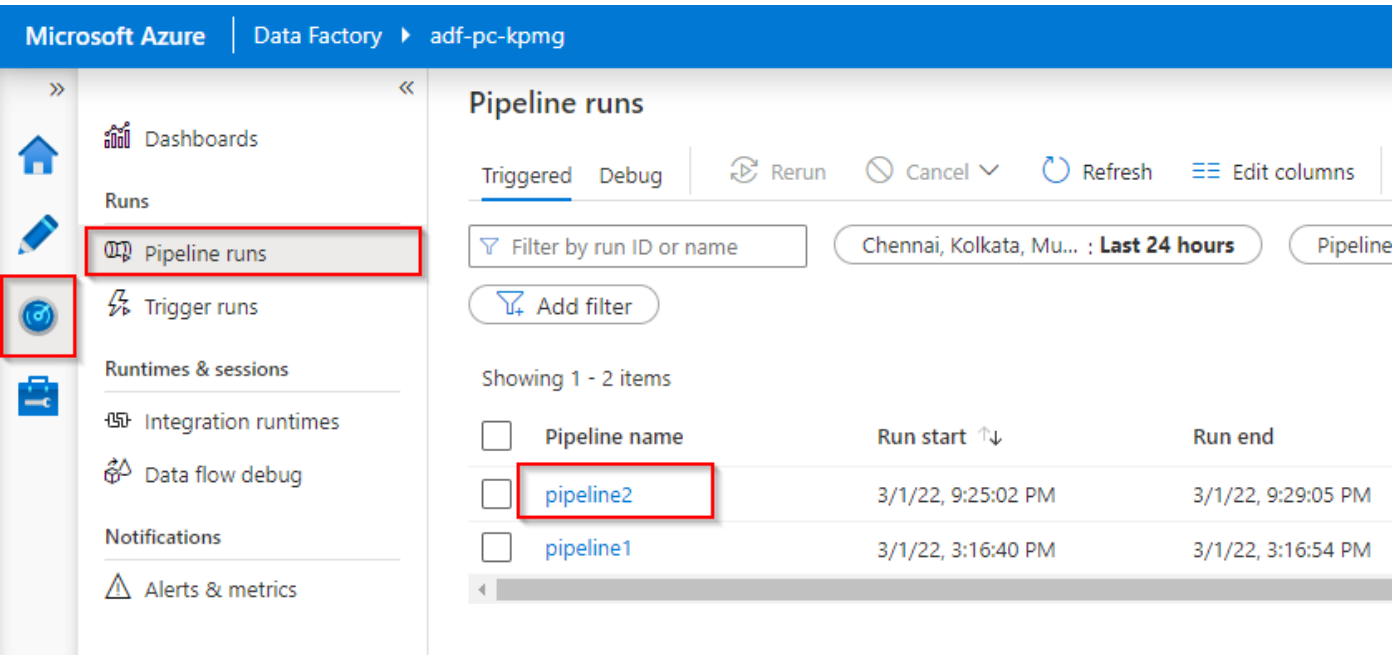


Task 5: Running the Pipeline

- 1. Go to the pipeline tab in the canvas and trigger the pipeline.



- 2. Go to monitor, under pipeline runs select the pipeline activity.



- 3. Select the pipeline activity run select eyeglass icon.

+
-
[100%]

Activity runs

Pipeline run ID 8b6200f6-5bf8-4583-8054-ef90f7551dc0

All status ▾

Showing 1 - 2 of 2 items

Activity name	Activity type	Run start ↑↓	Duration	Status	Error	Log
Data flow1	Data flow	3/1/22, 9:25:12 PM	00:03:52	✓ Succeeded		
Copy data1	Copy data	3/1/22, 9:25:03 PM	00:00:08	✓ Succeeded		

4. Monitor the status of each transformation.

✓ Data flow1

Cluster startup time: 3m 11s Number of transformations: 8 Data flow status: Success

Refresh

Auto refresh ☒ On ⓘ

Edit dataflow

- If you used the same logic described in this lab, your Data Flow should have written **737 rows** to your SQL Database.
- To verify this number navigate to you Azure SQL Database.
- In your SQL Database, navigate to the **Query Editor** under Tables click on ... and **select Top 1000 Rows** and verify the No. of rows in **Messages** Tab.

SQLDB (sqlservicepc/SQLDB) | Query editor (preview) ...

Search (Ctrl+/)

« Login + New Query ↕ Open query Feedback

- Overview
- Activity log
- Tags
- Diagnose and solve problems
- Quick start
- Query editor (preview)
- Power Platform
 - Power BI
 - Power Apps
 - Power Automate
- Settings
 - Compute + storage
 - Connection strings
 - Properties

SQLDB (sqladmin)

Showing limited object explorer here. For full capability please open SSDT.

- Tables
 - dbo.Ratings
- Views
- Stored Procedures

Query 1 × Query 2 ×

Run Cancel query Save query Export data as Show only Editor

1 SELECT TOP (1000) * FROM [dbo].[Ratings]

Results Messages

Affected rows: 737